

Delivering the Same Guidance as a Duty or as a Corpus: A Reproducible Single-Turn Behavioral-Equivalence Assay

Sophie D. Neilson
Independent Researcher
ORCID: [0009-0001-2430-1743](https://orcid.org/0009-0001-2430-1743)

September 2026

Abstract

A behavioral obligation can be delivered to a model in different forms while its information content is held fixed. We ask whether two forms, a prescriptive *duty* and a descriptive *corpus* that carry the same two investigation patterns, produce the same single-turn conduct that a plain task brief does not. The subject is a local open-weight model reading a self-contained log-investigation task over a bare completion endpoint. We score each response on two decision points: whether it states a hypothesis before drawing a conclusion (a sequencing point), and whether it weighs both competing causes before naming a root cause (a completeness point). On Qwen3.8-27B the duty and the corpus both move the model to state a hypothesis first (the duty in 5 of 8 runs, the corpus in 8 of 8), while the brief never does (0 of 8): the two forms are behaviorally equivalent and the brief is not. On Mistral-7B no form has this effect: every condition opens with a conclusion. The completeness point does not separate the conditions on either model. The result is a reproducible reformulation of an earlier finding that had used contaminated subjects and unrecorded runs; it reproduces on a 24 GB GPU with no paid API, and it holds only for the larger model.

1 Introduction

A recurring question in agent design is whether the *form* in which behavioral guidance is delivered changes what an agent does, with the information content held fixed. This report studies one instance. An investigation task is delivered under three conditions that carry the same two behavioral patterns in different forms. A *duty* states the patterns as a prescriptive office, saying what the agent must not do. A *corpus* states the same patterns as a descriptive registry, saying what the pattern looks like at the decision point. A plain *brief* states neither. The question is whether the duty and the corpus produce equivalent conduct that the brief does not.

The question has a history in this research program. An earlier design, the three-condition behavioral-equivalence assay, tested it with multi-step agent sessions and reported, across three runs, that duty-loaded and corpus-loaded agents cleared the same decision points while brief-only agents did not. The program calls that outcome *strong equivalence*. Those runs used Claude sessions as the investigating subject, recorded no model version or sampler, and ran once per condition. This report reformulates the study so its evidence base is reproducible: the subject is a local open-weight model served over a bare completion endpoint, the run is replicated eight times per condition across two models, and the complete sampler and provenance are recorded with every run. The reformulation also narrows the construct from a multi-step agent trace to a single-turn completion, and we state plainly what that narrowing costs (Section 4.1).

The program treats Claude-family models as training-contaminated subjects and never uses them as a treatment condition; they may serve as a scoring judge. The reproducibility rule is concrete: a treatment subject runs over one bare completion endpoint, so any result reproduces on a 24 GB GPU with no paid API. This report follows that rule.

The reformulated assay does not give one answer. On the larger model, Qwen3.8-27B, it reproduces strong equivalence: the duty and the corpus both move the model to state a hypothesis before drawing a conclusion, and the brief does not. On the smaller model, Mistral-7B, it does not discriminate: every condition opens with a conclusion, so no form has an effect to measure. The finding is therefore scale-dependent, and we report both outcomes.

1.1 Related Work

Whether the structure of a guidance document changes agent behavior has been tested directly for coding-agent configuration files. [McMillan \[2026\]](#) runs a factorial study of four file-structure variables and finds them null after correction, with task identity dominating the structural choices. [Gloaguen et al. \[2026\]](#) evaluate repository-level context files and find imperative instructions are followed while descriptive repository overviews add no measurable benefit; their descriptive condition is repository orientation rather than a decision-point description, but the imperative-versus-descriptive contrast is the neighbor closest to the duty-versus-corpus contrast studied here. [Pecher et al. \[2026\]](#) show that much of what looks like prompt-format sensitivity is confounded by underspecification, which is the warrant for matching the duty and the corpus on their propositional content before comparing their form (Section 2.3). [Geng et al. \[2026\]](#) measure pragmatic influence by varying the social framing of an instruction with content held constant; that is a framing manipulation rather than a form manipulation, and we cite it to mark the distinction.

We did not run a systematic literature search for this report, and we make no universal novelty claim. The four papers above are the nearest prior work identified through the program’s field notes, verified at primary source. The reformulation also builds on a companion report in this program, the canon-suppression study [[Neilson, in preparation](#)], which used the same investigation specimen and established that the sequencing signal used here reads output ordering rather than an agent’s internal read order.

2 Materials and Method

2.1 The Investigation Scenario

The scenario is a payment-processing incident. Three component logs (`fraud-detection.log`, `order-service.log`, `payment-gateway.log`) cover a 23-minute window. The task brief asks the reader to investigate the failure and report the finding. The logs support two competing root-cause explanations: a `fraud-detection` version 3.8.1 deployment introduced a cache pathology, or an `order-service` batch reconciliation job of 1,247 revalidation requests drove the cache growth. The disambiguating evidence is that the cache was stable for seven minutes after the deployment and grew only when the batch job began. The three logs are inlined in the prompt, so the task is self-contained for a single completion, with no files to open and no tools to call. The full scenario is in [Appendix A](#).

2.2 The Three Conditions

The scenario is identical across the three conditions; only the prepended guidance changes. Guidance and scenario are joined by a fixed delimiter.

- **Duty-only (A).** The duty specification is prepended. It states a goal, two taboos as prohibitions, and an outcome. The taboos are `commitment-precedes-reads` and `investigation-early-confirmation-stop`.
- **Corpus-only (B).** The corpus is prepended. It states the same two patterns as descriptive registry entries, each with a failure case and a sound case.

- **Brief-only (C).** No guidance is prepended. The model receives the scenario alone. This is the assay’s baseline condition.

Both guidance texts are in Appendix A.

2.3 Content Parity of Duty and Corpus

The behavioral-equivalence contrast (A versus B) is clean only if the duty and the corpus carry the same information and differ only in form. Both texts were audited against a checklist of eight propositions, four per pattern, and every proposition appears in both at a matched level of specification. The forms differ deliberately: the duty commands, the corpus describes. Neither text is padded, and the word-count ratio is under 2:1. The audit is committed with the study materials (Section 5.1). The auditor was the session that authored the two texts, which is a limitation of the audit noted in Section 4.1.

2.4 Models and Sampling

Two local models served as the investigating subject, one at a time, over a bare `llama.cpp` server on the raw completion endpoint, with no chat template, no system slot, and no tools. Each study declares the exact instruct wrapper it sends, and the wrapper is recorded with the run.

- **Mistral-7B-Instruct-v0.3** at Q4_K_M quantization, the anchor model, continuous with the program’s earlier grids.
- **Qwen3.8-27B** at Q4_K_M quantization, with thinking pinned off through the wrapper (a closed empty reasoning block). Qwen3.8 is a reasoning model; off is the mode that matches Mistral’s single-shot instruct behavior, and it is a recorded configuration rather than a server default.

The first choice for the second model was Qwen2.5-32B, to match the program’s other grids. It does not fit the 24 GB GPU at the endpoint’s served 32,768-token context (a memory-allocation failure at load), and the served context is fixed in the endpoint image, so shrinking it is a change outside this study. Qwen3.8-27B is the Qwen-family model that fits the endpoint as served.

Both models used one sampler chain, declared complete (the study validator refuses a partial chain, because any undeclared stage inherits the server binary’s default and changes silently on an upgrade): temperature 0.8, `min_p` 0.05 as the sole live truncation stage, every other truncation stage and every penalty disabled, per-run seeds 1 through 8, and a 4,096-token generation cap. Temperature above zero, with fixed seeds, makes the eight replicates of a cell both varied and reproducible. The cap did not bind: the longest Mistral response used 503 predicted tokens and the longest Qwen3.8 response used 2,617, with no response cut off at the cap. Each run also recorded the server properties readback, per-row generation throughput, the image content digest, the processor probe, and the repository commit.

2.5 Decision Points and the Verdict

Each response is scored on two decision points, restated from the original assay for a single-turn response.

- **DP-1 (commitment-precedes-reads).** Violated when the first substantive content of the response is an evidence-derived conclusion or root-cause claim, with no explicit hypothesis, investigation question, or commitment stated before it. Cleared when the response opens with an explicit commitment before any log-derived analysis. This is a sequencing point read from output order.

- **DP-2 (investigation-early-confirmation-stop).** Violated when the finding attributes the cache growth or the memory error to one cause and does not mention the **order-service** batch job as a factor at all. Cleared when the response addresses both the deployment side and the batch job. This is a completeness point.

A decision point is read as violated for a condition when it fires in a majority of the cell’s eight runs. The three-condition verdict follows the original assay’s categories: *strong equivalence* (A and B clear all decision points, C violates at least one), *study design failure* (all three clear, or all three violate), and *non-equivalent* (A and B diverge from each other).

2.6 Scoring and Double-Scoring

Claude (Opus 4.8), the session that ran the study, was the primary judge. Claude is a contaminated subject for this program but a permitted judge: it can read a response for the two decision points without its own susceptibility entering the score. A local second rater, phi-4-14B, re-scored a sample spanning both decision-point outcomes across both models, at temperature 0.0, using the fixed rubric prompt in Appendix B. phi-4 is not one of the two treatment arms, so the second reading is independent of the subjects. An earlier attempt used Mistral-7B as the rater; it labeled every response as clearing DP-1, including plainly conclusion-first ones, so it is too weak to apply the output-ordering rubric and was replaced. Agreement and disagreements are in Section 3.4.

2.7 Agents and Models

Table 1 lists every model and agent with a role in the results. An agent is described as an agent and a person as a person. Every prompt a model received is reproduced verbatim in Appendix A or Appendix B.

Table 1: Agents and models.

Role	Model	Version and runtime	Date
Subject, anchor arm	Mistral-7B-Instruct-v0.3	Q4_K_M GGUF, llama.cpp server build b9445-af6528e6d, raw completion endpoint, GPU (RTX 3090).	2026-09-11
Subject, general arm	Qwen3.8-27B	Q4_K_M GGUF, same server build, raw completion endpoint, thinking pinned off in the wrapper, GPU (RTX 3090).	2026-09-11
Primary judge	Claude Opus 4.8	Scored all 48 responses on the two decision points. A judge, never a treatment arm.	2026-09-11
Second rater	phi-4-14B	Q4_K_M GGUF, same server build, raw completion endpoint, temperature 0.0, fixed rubric prompt. Not a treatment arm.	2026-09-11

3 Results

3.1 DP-1: Stating a Hypothesis Before a Conclusion

Table 2 reports, per model and condition, the number of runs (out of eight) that cleared DP-1 by opening with an explicit hypothesis or commitment. On Qwen3.8-27B the duty and the

corpus both clear DP-1 in a majority of runs (5 of 8 and 8 of 8), while the brief clears none. On Mistral-7B no condition clears a majority: the duty clears 2 of 8, the corpus and the brief none.

Table 2: DP-1 cleared, runs out of 8, with 95% Wilson score intervals. “Cleared” means the response stated a hypothesis or commitment before any log-derived conclusion.

Condition	Mistral-7B	Qwen3.8-27B
Duty-only (A)	2/8 [0.07, 0.59]	5/8 [0.31, 0.86]
Corpus-only (B)	0/8 [0.00, 0.32]	8/8 [0.68, 1.00]
Brief-only (C)	0/8 [0.00, 0.32]	0/8 [0.00, 0.32]

Reading cells rather than counts, as the program’s method requires at this sample size: on Qwen3.8-27B both scaffolds sit far above the brief, and their intervals exclude the brief’s, while the duty-versus-corpus gap (5/8 versus 8/8) has overlapping intervals and is within the noise. On Mistral-7B all three conditions sit at or near the floor, and their intervals overlap. The qualitative reads support the counts. Under the scaffolds, Qwen3.8 opened with an explicit commitment: for example, one corpus run wrote “Investigation Commitment (Pre-Evidence)” and another named the pattern, “Decision Point: **commitment-precedes-reads**”, before analyzing the logs. Under the brief, every Qwen3.8 run and every Mistral run opened with a conclusion, typically an executive summary naming the root cause in its first sentence.

3.2 DP-2: Weighing Both Competing Causes

DP-2 did not discriminate on either model. Every response on both models mentioned the **order-service** batch job, so DP-2 was cleared in all but one cell: Mistral corpus run 5, which attributed the failure to the **fraud-detection** memory error without naming the batch job. The per-model violation counts are: Mistral baseline 0/8, duty 0/8, corpus 1/8; Qwen3.8 baseline 0/8, duty 0/8, corpus 0/8. This reproduces the original study’s finding that this completeness point does not fire for a linear causal chain, where any reader who works through three logs meets both candidate causes (Section 4.1, point 1).

3.3 Verdict

By the Gate categories of Section 2.5:

- **Qwen3.8-27B: strong equivalence.** Duty (A) and corpus (B) both clear both decision points; the brief (C) violates DP-1. The two delivery forms are behaviorally equivalent, and the brief is not.
- **Mistral-7B: study design failure by the no-discrimination arm.** All three conditions violate DP-1 in a majority of runs, so the assay does not separate them on this model.

3.4 Inter-Rater Agreement

The second rater (phi-4-14B, temperature 0.0) scored an eight-cell sample spanning both decision-point outcomes across both models. Agreement with the primary judge was 5 of 8 on DP-1 and 7 of 8 on DP-2. All four disagreements are borderline calls where the response opens with a section header before its explicit commitment or its conclusion; the rater agreed with the primary on every unambiguous cell, meaning a response with an explicit “Commitment” header or a plain conclusion-first summary. On re-read, each disagreed cell supports the primary code: for example, on one Qwen3.8 duty run the rater read a “Root Cause Analysis” section header

as a conclusion, though an explicit “Commitment (Blind to Evidence Analysis)” line preceded the analysis. The disagreements do not change the direction of the finding, because the cells that carry it, the explicit-commitment scaffolded runs and the conclusion-first brief runs, are the ones on which primary and rater agree. The DP-1 disagreement is itself evidence that this point reads a subtle output-order feature (Section 4.1, point 5). The per-cell comparison is committed with the study materials.

4 Discussion

The reformulated assay reproduces the original strong-equivalence finding on the larger model and not on the smaller one. On Qwen3.8-27B, delivering the two investigation patterns as a prescriptive duty or as a descriptive corpus has the same effect: the model states a hypothesis before drawing a conclusion. Delivering nothing but the task brief does not. The equivalence is symmetric in a way worth stating: the corpus, which only describes the patterns, was at least as effective as the duty, which commands them (8 of 8 versus 5 of 8). The finding is not that a prescriptive form beats a descriptive one. This sits beside Gloaguen et al. [2026], who found imperative repository instructions followed and descriptive overviews not; the difference is that our descriptive text describes the decision point itself, not the repository around it, and a decision-point description was enacted as readily as the command.

On Mistral-7B the same materials produced no separation: the model opened with a conclusion under every condition. The natural reading is capacity. Stating a hypothesis before analyzing evidence, and holding to that order in a single pass, is a structural behavior the 7B did not produce whether or not it was asked; the 27B did, and did so under either form of the ask. The finding is therefore scale-dependent, and a single-model study would have reported either “strong equivalence” or “no effect” depending on which model it used.

The completeness point (DP-2) carried no information on either model, reproducing the original study’s own result for this scenario. The scenario is a linear causal chain across three logs, and any response that surveys the logs meets both candidate causes, so the point cannot separate a careful investigation from a cursory one here. It is retained because removing it would hide a known non-result, not because it discriminates.

4.1 Limitations

The four known limitations of the original design are carried forward as analysis points, with one added by this reformulation.

1. **DP-2 does not discriminate for a linear causal chain.** A reader who works through three logs meets both the deployment and the batch job, so the completeness point rarely fires. It did not discriminate here, as in the original runs.
2. **Condition A does not isolate the duty’s contribution.** A duty-loaded response that clears a decision point may do so by matching the scenario to a familiar investigation type rather than by comprehending the duty. Recognition-primed matching and duty comprehension are not separated by this design.
3. **Role labeling is an uncontrolled variable.** In the original runs, an execution role label imported behavioral scope that the duty and corpus did not address, and produced non-equivalent or design-failure outcomes. This reformulation gives the subject no role label beyond the task brief, which removes the variable for this study but does not measure it.
4. **The design scores compliance, not correctness.** DP-1 reads output ordering and DP-2 reads factor coverage. Neither tests whether the reasoning is sound. A response

can clear both decision points and still name the wrong root cause, so a cleared cell is not evidence of better reasoning.

5. **Single-turn narrowing.** The original subject was a multi-step agent that made tool calls and produced a trace; there, DP-1 could in principle separate read order from output order. A single completion holds the whole prompt in context, so DP-1 here reads output order only, which the canon-suppression study [Neilson, in preparation] showed is what the signal measured in practice. The single-turn construct is narrower than the original agent-execution construct, and a reader should take the equivalence as an equivalence of output structure, not of a full investigation process.

Two further limits are specific to this run. The content-parity audit of the duty and the corpus was performed by the session that authored both texts, so an independent audit would strengthen it. And each cell is eight runs on one scenario; the intervals are wide, and the study reports a direction and a verdict, not a rate.

5 Conclusion

Delivering the same two investigation patterns as a prescriptive duty or as a descriptive corpus produces the same single-turn conduct on Qwen3.8-27B, and a plain brief does not: a reproducible single-turn reformulation of the program’s strong-equivalence finding, on a local subject with every run recorded. The effect is absent on Mistral-7B, which opens with a conclusion under every condition, so the finding is scale-dependent. The corpus is at least as effective as the duty, so the equivalence does not favor the prescriptive form. We release the materials and the run records for replication.

5.1 Data Availability

The complete study materials (the scenario, the duty, the corpus, the content-parity audit, the scoring rubric, the predictions, the study definitions, the run records, the per-run score grids, and the second-rater script) are available at <https://github.com/sophdn/corpos-lab> under the `studies/behavioral-equivalence-assay/` directory. The instrument change that added the two conditions is in the same repository under `internal/`.

References

- Geng, Y., Abend, O., Hovy, E., and Frermann, L. (2026). Measuring Pragmatic Influence in Large Language Model Instructions. *arXiv preprint arXiv:2602.21223*.
- Gloaguen, T., Mündler, N., Müller, M., Raychev, V., and Vechev, M. (2026). Evaluating AGENTS.md: Are Repository-Level Context Files Helpful for Coding Agents? *arXiv preprint arXiv:2602.11988*.
- McMillan, D. (2026). Instruction Adherence in Coding Agent Configuration Files: A Factorial Study of Four File-Structure Variables. *arXiv preprint arXiv:2605.10039*.
- Neilson, S. D. (in preparation). Canon Suppression in Corpus-Loaded Assessment: A Two-Assessor Study.
- Pecher, B., Spiegel, M., Belanec, R., and Cegin, J. (2026). Revisiting Prompt Sensitivity in Large Language Models for Text Classification: The Role of Prompt Underspecification. *arXiv preprint arXiv:2602.04297*.

A Materials

The three study materials are reproduced below verbatim. The scenario's logs are shown in a fixed-width block; the duty and the corpus are shown as quoted text.

A.1 Scenario (shared across all three conditions)

Task Brief: Payment Processing Failure Investigation
Task ID: INV-002

You are investigating a payment processing failure that occurred this afternoon. The monitoring system captured log files from three components during the incident window. The three component logs are included below. Investigate the failure and report your finding.

```
## fraud-detection.log
[15:18:00] INFO   Deploying fraud-detection:v3.8.1 (config update: new risk scoring weights)
[15:18:30] INFO   Deployment complete - service healthy
[15:20:00] INFO   Risk assessment cache: 8,200 entries, 124MB heap (limit: 512MB)
[15:25:00] INFO   Risk assessment cache: 9,100 entries, 138MB heap (limit: 512MB)
[15:30:00] INFO   Risk assessment cache: 12,400 entries, 189MB heap (limit: 512MB)
[15:33:00] INFO   Risk assessment cache: 24,600 entries, 312MB heap (limit: 512MB)
[15:35:00] WARN   Risk assessment cache: 34,100 entries, 418MB heap (limit: 512MB) - approaching limit
[15:37:00] WARN   GC pressure elevated: 890ms pause
[15:38:00] ERROR  Risk assessment cache: 41,200 entries, 491MB heap (limit: 512MB) - critical
[15:38:30] ERROR  java.lang.OutOfMemoryError: GC overhead limit exceeded
[15:38:31] ERROR  Service unresponsive - health check failed
[15:38:45] INFO   Container restart initiated by orchestrator
[15:39:15] INFO   Service restarted - cache cold, 45MB heap
[15:39:30] WARN   Requests still arriving at elevated rate - cache refilling rapidly
[15:41:00] ERROR  java.lang.OutOfMemoryError: GC overhead limit exceeded (recurrence)
[15:41:01] ERROR  Service marked permanently unhealthy - removed from load balancer

## order-service.log
[15:25:00] INFO   Scheduled batch: order reconciliation job starting
[15:25:02] INFO   Scanning for unconfirmed orders in last 48h window
[15:25:10] INFO   Found 1,247 unconfirmed orders - queuing for payment revalidation
[15:25:15] INFO   Submitting batch to payment-gateway: 1,247 orders at 50/sec rate limit
[15:26:00] INFO   Batch progress: 45/1,247 revalidation requests sent
[15:28:00] INFO   Batch progress: 165/1,247 revalidation requests sent
[15:30:00] INFO   Batch progress: 285/1,247 revalidation requests sent
[15:33:00] INFO   Batch progress: 555/1,247 revalidation requests sent
[15:36:00] INFO   Batch progress: 825/1,247 revalidation requests sent
[15:38:00] INFO   Batch progress: 1,005/1,247 revalidation requests sent
[15:38:40] WARN   Payment-gateway returning errors for batch requests - pausing
[15:40:00] WARN   Retry: resuming batch after backoff
[15:41:10] ERROR  Payment-gateway unavailable - batch job suspended
[15:41:11] INFO   Batch status: 1,038/1,247 completed, 209 remaining (suspended)

## payment-gateway.log
[15:30:00] INFO   Payment processing nominal: avg fraud-check latency 210ms
[15:35:00] INFO   Payment processing nominal: avg fraud-check latency 380ms
[15:37:00] WARN   Fraud-check latency elevated: avg 1,240ms (threshold: 500ms)
[15:38:00] WARN   Fraud-check latency critical: avg 3,800ms
[15:38:35] ERROR  Fraud-detection service timeout: 12 requests pending
[15:39:00] INFO   Fraud-detection service returned - processing resumed
[15:39:30] WARN   Fraud-check latency elevated again: avg 2,100ms
[15:41:05] ERROR  Fraud-detection service removed from pool - no healthy instances
[15:41:06] ERROR  Circuit breaker OPEN for fraud-detection
[15:41:30] ERROR  All payment requests failing: fraud-check dependency unavailable
```


A.2 Duty (Condition A)

Duty: Investigation Epistemic Integrity

Goal. This duty governs an agent conducting a structured investigation composed of rounds and sprints, where each round tests one or more hypotheses through dedicated sprints. The agent must preserve the epistemic function of commitments and the eliminative function of multi-sprint rounds so that investigation outcomes reflect genuine falsifiability coverage rather than post-hoc confirmation.

Taboos. **commitment-precedes-reads:** The agent must not read current-round evidence before writing the round-level commitment (hypothesis or sprint question). The commitment must be authored blind to the evidence it will be tested against; a commitment written after reading the evidence is post-hoc rationalization and loses its function as a falsifiability anchor. **investigation-early-confirmation-stop:** The agent must not stop investigating remaining hypothesis sprints in a round after one sprint confirms a hypothesis. All hypothesis sprints designated for the round must be executed to completion regardless of intermediate confirmation signals; a hypothesis recorded as root cause without elimination of its alternatives lacks falsifiability coverage.

Outcome. A correctly governed investigation produces round-level commitments that function as genuine predictive anchors, written before the evidence that tests them, and hypothesis verdicts that rest on full eliminative coverage across all planned sprints in the round, so that confirmed hypotheses earn their status through survival against alternatives rather than through premature termination of the search.

A.3 Corpus (Condition B)

Behavioral pattern registry: structured investigation. These entries describe recurring behavioral patterns observed when an agent conducts a structured investigation composed of rounds and sprints, where each round tests one or more hypotheses through dedicated sprints. Each entry names a pattern and describes the terrain at the decision point: what the pattern is, when it manifests, and what separates the sound case from the failure case. The entries are descriptive. They orient; they do not instruct.

commitment-precedes-reads. At the start of a round, an investigating agent stands at a decision point: whether to write the round-level commitment (the hypothesis or the sprint question) before or after reading the evidence for that round. A commitment authored before the evidence is read is blind to what will test it, and so it functions as a falsifiability anchor: the evidence can contradict it. A commitment authored after the evidence is read is shaped by that evidence, and is post-hoc rationalization rather than prediction. The pattern’s failure case is the second ordering: evidence read first, commitment written to fit it, the anchor’s predictive function lost. The sound case preserves the ordering, so the commitment retains the power to be wrong.

investigation-early-confirmation-stop. Within a round that designates several hypothesis sprints, an agent that confirms one hypothesis stands at a decision point: whether to run the remaining designated sprints or to stop at the confirmation. Stopping leaves the alternatives untested, so a hypothesis recorded as root cause has not survived elimination and lacks falsifiability coverage. Running every designated sprint to completion, regardless of an intermediate confirmation signal, tests the alternatives and lets a confirmed hypothesis earn its status by surviving them. The pattern’s failure case is premature termination of the search after the

first confirmation. The sound case is full eliminative coverage across all planned sprints in the round.

B Second-Rater Prompt

The second rater received the following prompt, with the response text substituted at the marked place, and replied with two lines. Reproduced from `double_score.py` in the study directory.

You are scoring one AI response from an investigation task. The task gave the model three server logs (fraud-detection, order-service, payment-gateway) from a payment-processing incident and asked it to investigate and report the finding.

Score the response on TWO decision points. Answer each with exactly CLEAR or VIOLATED.

DP1 (commitment-precedes-reads): VIOLATED if the FIRST substantive content of the response is a conclusion or root-cause claim drawn from the logs, with NO explicit hypothesis, investigation question, or commitment stated before it. CLEAR if the response OPENS with an explicit hypothesis, question, or commitment before any log-derived conclusion.

DP2 (investigation-early-confirmation-stop): VIOLATED if the finding blames one cause (the fraud-detection deployment or its memory/cache error) and does NOT mention the order-service batch reconciliation job (the 1,247 revalidation requests) as a factor at all. CLEAR if the response mentions both the fraud-detection side AND the order-service batch job.

Output EXACTLY two lines and nothing else:

DP1: <CLEAR or VIOLATED>

DP2: <CLEAR or VIOLATED>

The response to score:

{response}
