

Canon Suppression in Corpus-Loaded Assessment: A Two-Assessor Study

Sophie D. Neilson
Independent Researcher
ORCID: [0009-0001-2430-1743](https://orcid.org/0009-0001-2430-1743)

September 2026

Abstract

We report a study comparing the assessments of a corpus-loaded evaluator and a corpus-naive observer applied to the same behavioral data. A canon-carrying Researcher assessed an agent execution trace as part of an ongoing research program; a vanilla Claude session—carrying no role, no corpus, and no assessment vocabulary—independently reviewed the same trace. The vanilla observer flagged 7 items as significant. Of these, the corpus-loaded Researcher had absorbed 3 (data present but analyzed into non-significance) and was blind to 4 (data never reached attention). Zero items were correctly excluded. Two distinct suppression mechanisms were identified: *verdict-space limitation* (the assessment instrument has no slot for the finding) and *canon-frame limitation* (the assessor’s corpus knowledge makes the finding appear expected or non-significant). The strongest single datum—an agent producing correctly-structured eliminative reasoning while reaching an unsupported conclusion—was a pure canon-frame suppression: the instrument could have caught it, but the canon made the form of good reasoning indistinguishable from its substance. We discuss the finding as an instance of the observer-position problem in evaluation methodology, drawing on Bourdieu’s reflexive response and distinguishing it from Garfinkel’s naive-observer fix.

1 Introduction

When an evaluator is deeply familiar with the system under study, that familiarity creates a structural blind spot. The evaluator’s knowledge provides a template for what correct behavior looks like, and that template can make genuine findings appear expected—hence non-significant—before they reach the verdict.

This is not motivational bias. The evaluator is not defending the system. The evaluator’s expertise makes certain data features *invisible*: they look like what correct behavior should produce, so they do not register as findings. The result is that expertise, which is required for competent evaluation, simultaneously suppresses the evaluator’s ability to notice certain classes of evidence.

We encountered this problem concretely during a behavioral equivalence study on language model agents [Neilson, 2026a]. A Researcher who had designed the study environment, written the evaluation instrument, and maintained the behavioral corpus assessed a set of agent execution traces. We then ran the same traces through a vanilla Claude session carrying no corpus, no role definition, and no assessment vocabulary. The vanilla observer flagged 7 items as significant. The Researcher had suppressed all 7.

This paper reports that study, classifies the suppression mechanisms, and frames the result within the observer-position literature. The contribution is not the existence of evaluator bias—that is well established. The contribution is the *mechanism taxonomy*: two structurally distinct suppression types with different design responses, identified in a concrete case within AI behavioral evaluation.

2 Study Design

2.1 Specimen

The specimen was Run 6 of a behavioral equivalence assay: three agent execution traces produced under different specification-loading conditions (Conditions A, B, and C), assessed against a signal record with pre-defined decision points. The task scenario was a simulated production incident requiring root-cause analysis. The environment was designed such that two competitive explanations existed: a deployment configuration change and a batch reconciliation job. The environment designer (the same Researcher who later assessed the traces) considered the batch job the stronger explanation.

2.2 Two-Assessor Design

Assessor 1 (Researcher). A canon-carrying evaluator: the person who designed the study environment, authored the evaluation instrument, maintained the behavioral corpus, and held the full context of the research program. The Researcher assessed the traces using the pre-defined signal record and produced a verdict with supporting analysis.

Assessor 2 (Vanilla observer). A Claude session with no role file, no corpus access, no registry, no taboo vocabulary, and no prior run context. The vanilla observer received the three execution traces, the task brief, the environment logs, and the Researcher’s completed study record. It was prompted to identify anything in the traces that the study record did not flag as significant.

The design is deliberately asymmetric. Assessor 1 evaluates the traces under normal conditions. Assessor 2 reads both the traces and the evaluation, looking for what the evaluation missed. The comparison surfaces what the canon suppressed.

3 Results

The vanilla observer flagged 7 items. The Researcher annotated each item with its suppression status and mechanism.

3.1 The Seven Items

V1—Root cause divergence. Conditions A and B concluded the deployment was the root cause and dismissed the batch job. Condition C inverted this: the batch job was the root cause, the deployment a contributing factor. The Researcher had noted this divergence but classified it as outside the signal record’s scope. *Status: absorbed.*

V2—Factual error in dismissing the batch job. Condition A argued that cache growth preceded the batch job’s throughput, citing a figure (285 requests by 15:30) that referred to the period *before* the acceleration. The actual acceleration occurred while the batch job was at sustained throughput. The Researcher’s own environment design described the batch job as the trigger, but never verified the agents’ evidence citations against the log data. *Status: invisible.*

V3—Log-before-commitment in Condition B. Condition B admitted reading logs before forming hypotheses, but framed its output as compliant (hypotheses before analysis in layout). The Researcher flagged this problem in Condition A, which self-disclosed it as a violation, but not in Condition B, which framed it as compliance. *Status: absorbed.*

V4—Condition C produced richer practical output. Condition C included recommendations, impact assessment, and system design gap identification—none present in A or B. The Researcher noticed this but the instrument measures behavioral compliance, not output quality. *Status: absorbed.*

V5—Conditions A and B converge on content, not just structure. Same hypotheses, same dismissal framing, same remediation. The specification channeled the analytical pathway, not just the output format. The Researcher did not notice the depth of convergence. *Status: invisible.*

V6—Environment design contradicts scoring. The Researcher’s environment rationale (batch job is the competitive explanation) did not carry into the assessment (agents’ dismissal of the batch job accepted without cross-reference). *Status: invisible.*

V7—Condition C’s incident window. Condition C placed the incident origin at 15:25 (batch job start), not 15:18 (deployment). A direct expression of the different causal attribution, not noted by the Researcher. *Status: invisible.*

3.2 Classification

Table 1: Suppression classification of the seven flagged items.

Item	Verdict-space	Canon-frame	Primary mechanism
V1	Yes	Yes	Both
V2	No	Yes	Canon-frame only
V3	Yes	Yes	Both
V4	Yes	Partial	Verdict-space
V5	Yes	Yes	Both
V6	Yes	Yes	Both
V7	Yes	Mild	Verdict-space

Three items were absorbed (data present, analyzed away). Four were invisible (data never reached attention). Zero were correctly excluded—every item the vanilla observer flagged was either real signal the Researcher had suppressed, or real signal the Researcher had never seen.

4 Two Suppression Mechanisms

4.1 Verdict-Space Limitation

The assessment instrument cannot express certain findings. The signal record defines decision points and verdict codes; data features that fall outside these definitions have no slot. Items V4 (output quality), V5 (content convergence), V6 (design-assessment contradiction), and V7 (framing metadata) were invisible or absorbed partly because the instrument provided no way to record them.

The design response to verdict-space limitation is *instrument revision*: add slots for the missing finding types.

4.2 Canon-Frame Limitation

The assessor’s corpus knowledge makes a finding appear expected. The specification describes what correct behavior looks like; when the agents produce output that matches the expected form, the evaluator does not inspect whether the substance matches the form. Items V1, V2, V3, V5, and V6 were suppressed partly because the canon provided a template that the agents’ output satisfied formally.

The design response to canon-frame limitation is *position declaration*: make the assessor’s epistemic position legible as part of the research record, so that readers can distinguish position-dependent judgments from position-independent ones.

4.3 Reinforcement

In 4 of 7 cases, both mechanisms reinforced each other: the instrument had no slot for the finding *and* the canon made it look expected. Two items were primarily verdict-space. One item—V2—was purely canon-frame.

5 V2: The Critical Datum

V2 merits separate treatment because it is a pure canon-frame suppression in which the instrument *could* have caught the finding.

The specification-loaded agents (Conditions A and B) dismissed the batch job as the root cause. They did so using structured eliminative reasoning: separate investigation sprints per hypothesis, a verdict per sprint, evidence citations. The form of the reasoning was correct. The substance was not—the evidence citations did not support the dismissal.

The Researcher never verified this. The structured elimination output looked like sound elimination to a canon-carrying assessor. The canon provides a template for what good investigation looks like (structured elimination with evidence citations), and the agents’ output satisfied that template. The Researcher saw the form and inferred the substance.

This is not carelessness. It is what the canon does: it provides a recognition template, and that template can be satisfied formally while the epistemic function—actually checking whether the evidence supports the conclusion—is not served. V2 demonstrates that a canon-carrying assessor can look at correctly-structured reasoning and miss that the reasoning is wrong.

5.1 Implications for Formal versus Substantive Compliance

V2 bears directly on the distinction between formal and substantive compliance in language model behavior. If a specification produces agents that follow the *form* of eliminative investigation while reaching a conclusion unsupported by the data, what we observe is formal compliance without substantive comprehension. The agents learned what elimination looks like—separate sprints, per-sprint verdicts—without what elimination does—checking whether the evidence actually supports the dismissal.

This is constraining evidence for the Comprehension-as-Compliance hypothesis [Neilson, in preparation], which claims that structured descriptions of decision classes produce genuine compliance through recognition. V2 suggests the claim may hold for behavioral pattern compliance (correct navigation at decision points) but not for epistemic-function compliance (actually performing the reasoning the patterns encode). One specimen, one environment: the constraint is real but the scope is narrow.

6 The Observer-Position Problem

The canon-suppression finding is an instance of the observer-position problem: the evaluator’s position within the research apparatus affects what they can see.

6.1 Why Garfinkel’s Fix Does Not Transfer

Garfinkel’s ethnomethodological response to self-report bias is to replace the insider with a naive observer—someone who lacks the cultural competence that produces shaped accounts [Garfinkel, 1967]. The vanilla observer in this study plays exactly that role, and its novelty detection worked: 7/7 items were genuine signal.

But the Garfinkelian fix is not available at the assessment level. A naive observer lacks the corpus competence the assessment *requires*. The vanilla observer could flag “this looks surprising” but could not evaluate whether the agents’ behavior was correct *relative to the*

specification. Naive observation and competent assessment are structurally incompatible: the competence that enables the assessment is the competence that produces the suppression.

6.2 Bourdieu’s Reflexive Response

Bourdieu’s response is different: rather than replacing the biased observer, the observer’s position is made legible [Bourdieu and Wacquant, 1992]. The objectification of the objectifying subject—declaring the assessor’s epistemic position as part of the research apparatus—does not eliminate the bias. It makes the bias *readable*. A reader who knows the assessor’s position can distinguish “the Researcher found this unremarkable” (a position-dependent judgment) from “this is unremarkable” (a position-independent claim).

The canon-suppression study is itself the position declaration. It makes “the Researcher found correct navigation obvious” into an observable artifact rather than an invisible assumption.

7 Limitations

1. **One specimen.** The study assessed one set of execution traces from one study environment. The suppression patterns may not generalize.
2. **One assessor.** A single Researcher produced both the environment and the assessment. The suppression may be specific to this individual’s relationship with this corpus.
3. **Uncalibrated vanilla observer.** The vanilla observer flagged 7 items and all 7 were genuine. This does not mean vanilla observers have zero false positives in general—the 7/7 rate may reflect this particular trace set rather than a stable property.
4. **Root-cause ambiguity.** V2’s status as the “critical datum” depends on the claim that the agents reached the wrong root cause. The evidence is genuinely ambiguous: the deployment configuration change plausibly altered cache behavior independently of the batch job. The environment designer believed the batch job was the trigger, but that belief is itself a canon-carrying judgment.
5. **Log-analysis domain only.** All traces involved incident investigation. The suppression mechanisms may operate differently in other task domains.

8 Conclusion

A corpus-loaded evaluator assessing a corpus-designed study environment suppressed 7 out of 7 items flagged by a corpus-naive observer. Zero items were correctly excluded. The suppression operates through two distinct mechanisms: verdict-space limitation (the instrument cannot express the finding) and canon-frame limitation (the assessor’s knowledge makes the finding appear expected). The two mechanisms have different design responses: instrument revision and position declaration, respectively. The strongest single finding is a case where an agent produced correctly-structured reasoning while reaching an unsupported conclusion—a pure canon-frame suppression where the form of good reasoning masked the absence of its substance.

The result is an instance of the observer-position problem, and the Bourdieu-reflexive response—making the assessor’s position legible—is the structural fix. The study itself is that fix: it converts “the evaluator found this unremarkable” from an invisible assumption into an observable artifact.

8.1 Data Availability

The complete study record, including the vanilla observer’s findings, the Researcher’s annotations, and the suppression mechanism classification, is available at <https://github.com/sophdn/corpos-lab>.

References

- Bourdieu, P. and Wacquant, L. J. D. (1992). *An Invitation to Reflexive Sociology*. University of Chicago Press.
- Garfinkel, H. (1967). *Studies in Ethnomethodology*. Prentice-Hall.
- Neilson, S. D. (2026). Structured phenomenological descriptions induce analysis-mode behavior in a small language model. Zenodo. [doi:10.5281/zenodo.22542747](https://doi.org/10.5281/zenodo.22542747).
- Neilson, S. D. (in preparation). Comprehension-as-Compliance: Structured phenomenological descriptions as a steering mechanism for language model agents.